

Knowledge Boundary Identification

Moxin Li

National University of Singapore

Speaker Info



Homepage



I am a fourth-year Ph.D. candidate at NExT Research Center, National University of Singapore (NUS).

Current research interests: trustworthy LLM, LLM safety, LLM evaluation.

Experience:

2021 - Now: Ph.D, NUS, School of Computing

2016 - 2020: B.S., Peking University, Yuanpei College

Email: limoxin@u.nus.edu

(**Actively looking for postdoc positions!**)



Knowledge Boundary Identification

Outline

❑ **Uncertainty Estimation**

- ❑ Quantifies the model's uncertainty about predictions.
- ❑ High uncertainty tends to lie outside of boundary.

❑ **Confidence Calibration**

- ❑ Aligns LLM's confidence with actual correctness of predictions.
- ❑ High confidence tends to lie inside of boundary.

❑ **Internal State Probing**

- ❑ Probes LLM internal states (e.g., attention heads, hidden layers, neurons) to assess factual accuracy.



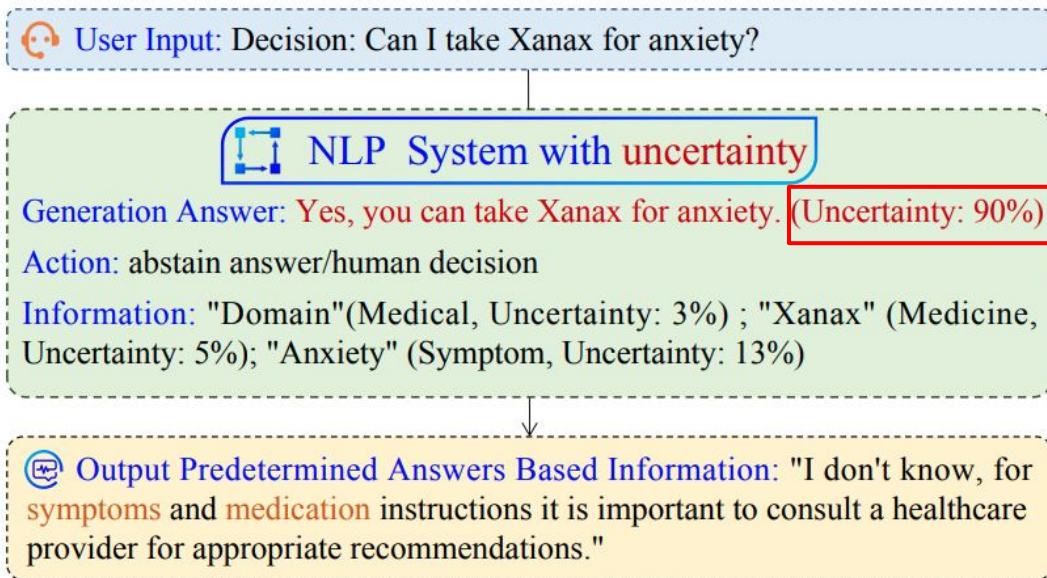
Knowledge Boundary Identification

Outline

- ❑ Uncertainty Estimation
 - ❑ Token Probability-based
 - ❑ Semantic-based
 - ❑ Uncertainty Decomposition
 - ❑ Conformal Prediction
- ❑ Confidence Calibration
- ❑ Internal State Probing

Uncertainty Estimation

LLM, as a neural network, makes mistakes.
Estimating the reliability of LLM's output is important.



Uncertainty Estimation (UE) – Token Probability-based

Question Answering

Prompt


How long did the
Roman Empire last?



LLM

... Lasted for over 500 years

Token Probability



Max/Avg Prob.

$$\text{Max}(-\log p)_i = \max_j (-\log p_{ij}),$$

$$\text{Avg}(-\log p)_i = -\frac{1}{J} \sum_j \log p_{ij},$$

Max/Avg Entropy

$$\text{Max}(\mathcal{H})_i = \max_j [\mathcal{H}_{ij}],$$

$$\text{Avg}(\mathcal{H})_i = \frac{1}{J} \sum_j \mathcal{H}_{ij},$$

Uncertainty Estimation (UE) – Token Probability-based

Different token weights and granularity

Question: What is the ratio of the mass of an object to its volume?


Ground Truth: density






LLMs Generation: density of an object



Predictive Entropy-based Uncertainty Quantification



Shifting Attention to Relevance Uncertainty Quantification



Uncertainty Estimation (UE) – Semantic-based

Semantic Entropy

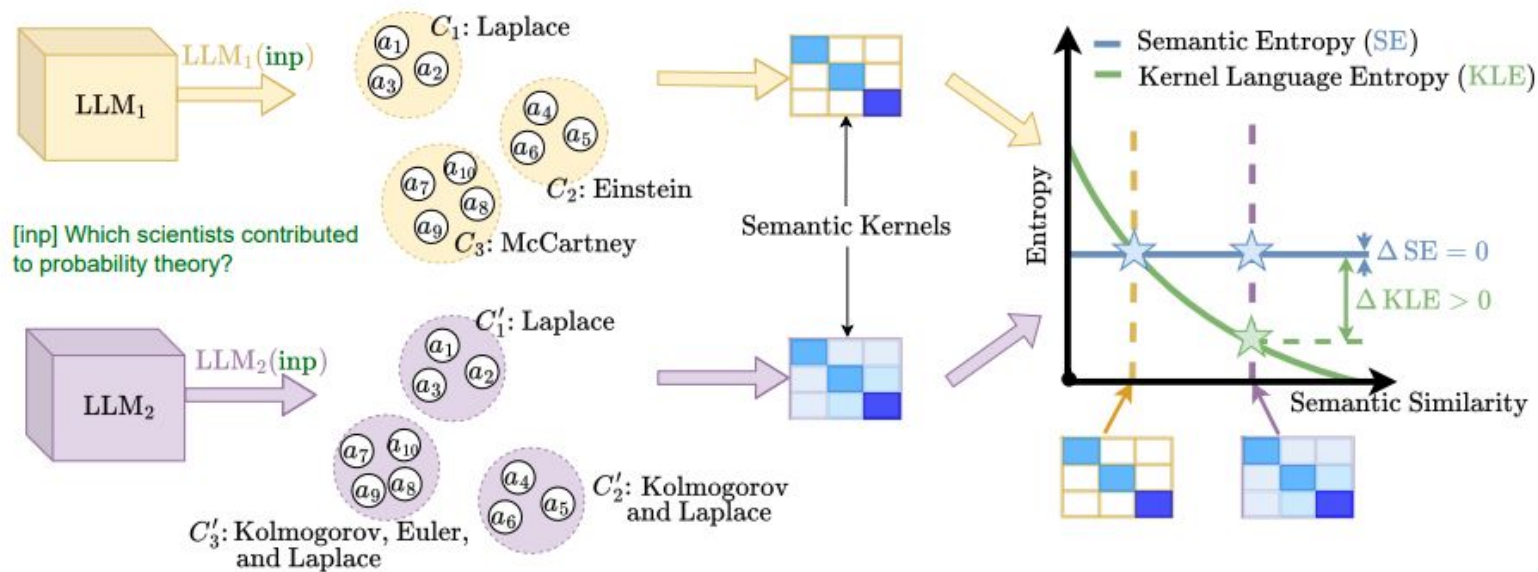
1. **Generation:** Sample M sequences $\{s^{(1)}, \dots, s^{(M)}\}$ from the predictive distribution of a large language model given a context x .
2. **Clustering:** Cluster the sequences which mean the same thing using our bi-directional entailment algorithm.
3. **Entropy estimation:** Approximate semantic entropy by summing probabilities that share a meaning following Eq. (2) and compute resulting entropy. This is illustrated in Table 1.

$$p(c | x) = \sum_{\mathbf{s} \in c} p(\mathbf{s} | x) = \sum_{\mathbf{s} \in c} \prod_i p(s_i | s_{<i}, x).$$

(a) Scenario 1: No semantic equivalence			(b) Scenario 2: Some semantic equivalence		
Answer \mathbf{s}	Likelihood $p(\mathbf{s} x)$	Semantic likelihood $\sum_{\mathbf{s} \in c} p(\mathbf{s} x)$	Answer \mathbf{s}	Likelihood $p(\mathbf{s} x)$	Semantic likelihood $\sum_{\mathbf{s} \in c} p(\mathbf{s} x)$
Paris	0.5	0.5	Paris	0.5	} 0.9
Rome	0.4	0.4	It's Paris	0.4	
London	0.1	0.1	London	0.1	
Entropy	0.94	0.94	Entropy	0.94	0.33

Uncertainty Estimation (UE) – Semantic-based

Kernel Language Entropy: considering inter-cluster similarity



Uncertainty Estimation – Uncertainty Decomposition

Uncertainty

=

Epistemic
Uncertainty

+

Aleatoric
Uncertainty

Model uncertainty:

- Model lack of knowledge
- Suboptimal modeling
- Perturbation randomness
- Reducible with stronger model

~ **Parametric Knowledge Boundary**
- **Outward Knowledge Boundary**

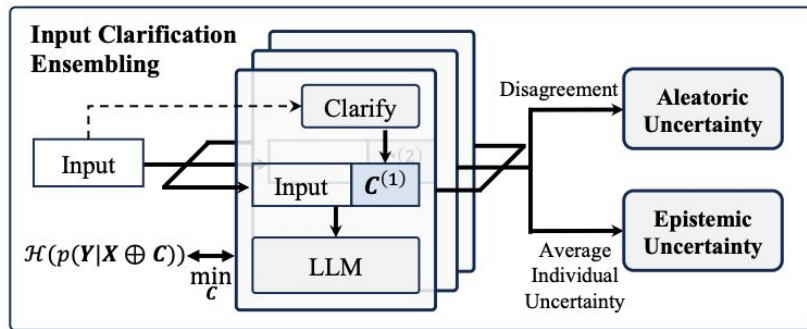
Data uncertainty:

- Question ambiguity
- Multiple answers
- Generation randomness (entropy)
- Less likely to be reducible

~ **Outward Knowledge boundary**

Most current works **do not distinguish** the two types of uncertainty and focus on the general identification of the **Outward Knowledge Boundary**.

Uncertainty Estimation – Uncertainty Decomposition



$$\mathcal{H}(q(Y|X)) = \underbrace{\mathcal{I}(Y; C|X)}_{\textcircled{1}'} + \underbrace{\mathbb{E}_{q(C|X)} \mathcal{H}(q(Y|X \oplus C))}_{\textcircled{2}'}.$$

Aleatoric uncertainty (1):

mutual information between the model output distribution and the clarifications.

Epistemic uncertainty (2): average entropy of the output distribution given different clarifications.

Uncertainty Estimation – Conformal Prediction

For classification task, conformal prediction produces a prediction set of labels $\mathcal{C}(X_t) \subset \mathcal{Y}$

$$p(Y_t \in \mathcal{C}(X_t)) \geq 1 - \alpha,$$

1. Identify a heuristic notion of uncertainty based on the model f ;
2. Define a conformal score function $s(X, Y) \in \mathbb{R}$ with larger scores encoding worse agreement between X and Y ; **e.g., using softmax score corresponding to the true label** $s(X, Y) = 1 - f(X)_Y$
3. Compute conformal scores on the calibration set $s_1 = s(X_c^{(1)}, Y_c^{(1)}), \dots, s_n = s(X_c^{(n)}, Y_c^{(n)})$ and calculate a threshold \hat{q} as the $\frac{\lceil (n+1)(1-\alpha) \rceil}{n}$ quantile of the calibration scores,

$$\hat{q} = \text{quant}\left(\{s_1, \dots, s_n\}, \frac{\lceil (n+1)(1-\alpha) \rceil}{n}\right), \quad (2)$$

where $\lceil \cdot \rceil$ is the ceiling function;

4. Construct the prediction set for each test instance X_t as

$$\mathcal{C}(X_t) = \{Y' \in \mathcal{Y} : s(X_t, Y') \leq \hat{q}\}. \quad (3)$$

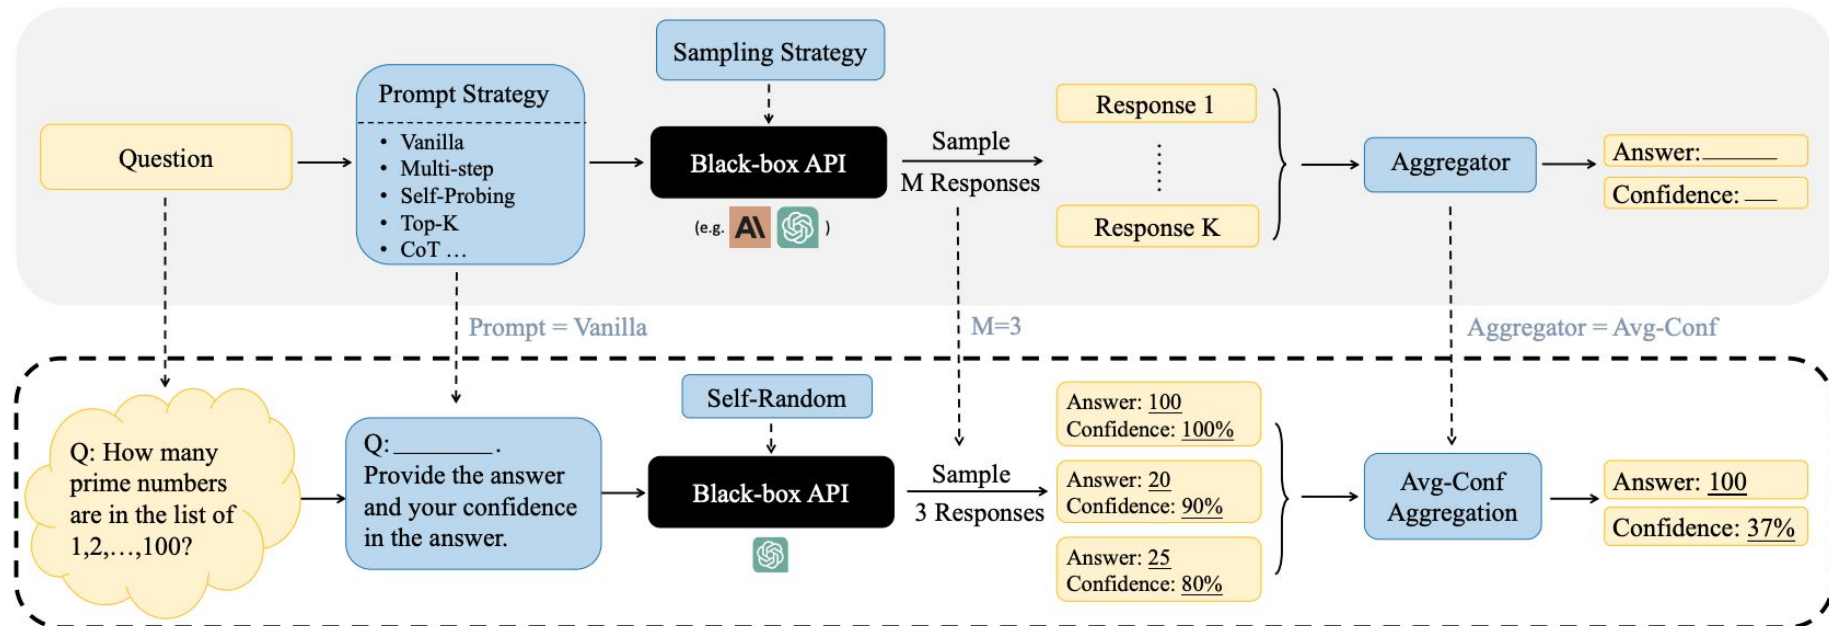


Knowledge Boundary Identification

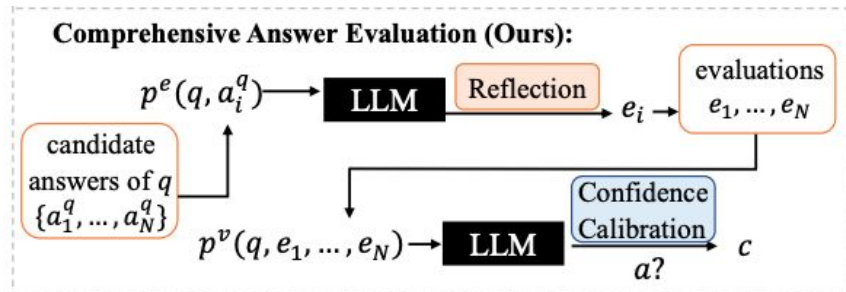
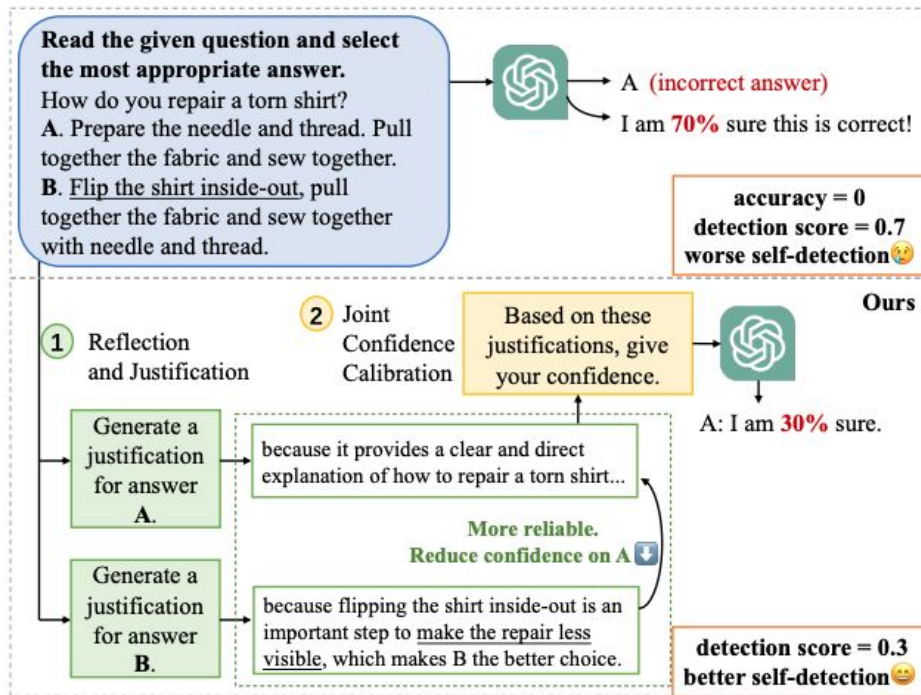
Outline

- ❑ Uncertainty Estimation (UE)
- ❑ Confidence Calibration
 - ❑ Prompt-based Calibration
 - ❑ Fine-tuning for Calibration
- ❑ Internal States Probing

Confidence Calibration – Prompt-based Calibration



Confidence Calibration – Prompt-based Calibration



Confidence Calibration – Prompt-based Calibration

$P(\text{True})$

Question: Who was the first president of the United States?



Proposed Answer: George Washington

Is the proposed answer:

(A) True

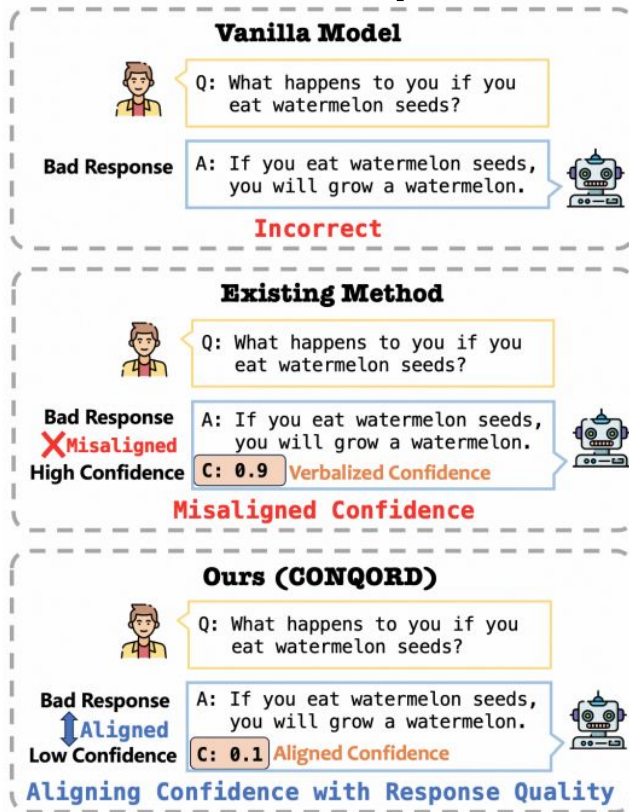
(B) False

The proposed answer is:



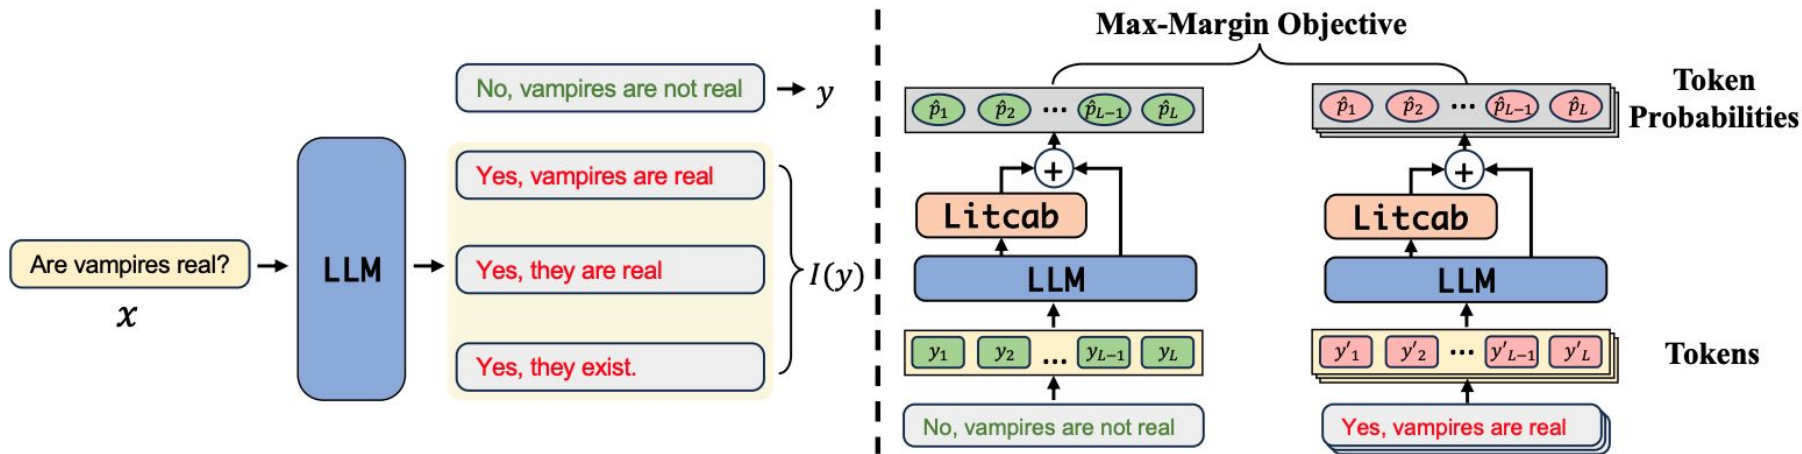
Confidence Calibration – Fine-tuning for Calibration

Fine-tuning for verbalized confidence expression



Confidence Calibration – Fine-tuning for Calibration

Adjusting token probability





Knowledge Boundary Identification

Outline

- ❑ Uncertainty Estimation (UE)
- ❑ Confidence Calibration
- ❑ Internal State Probing

Internal State Probing – Inference-Time Intervention (ITI)

Motivation

LLMs often know factual truth internally but still output falsehoods.

Key Observation

There's ~40% gap between what LLMs' hidden activations encode (via probe accuracy) vs. what they generate (output accuracy).

Proposed Method (ITI)

- During inference, identify a sparse subset of attention heads whose activations correlate with truth (measured by TruthfulQA).
- Shift activations along these "truthful directions" to nudge output toward truth.

Result

On Alpaca (instruction-tuned LLaMA), truthfulness jumped from 32.5% → 65.1% on TruthfulQA.

Internal State Probing – Universal Truthfulness Hyperplane

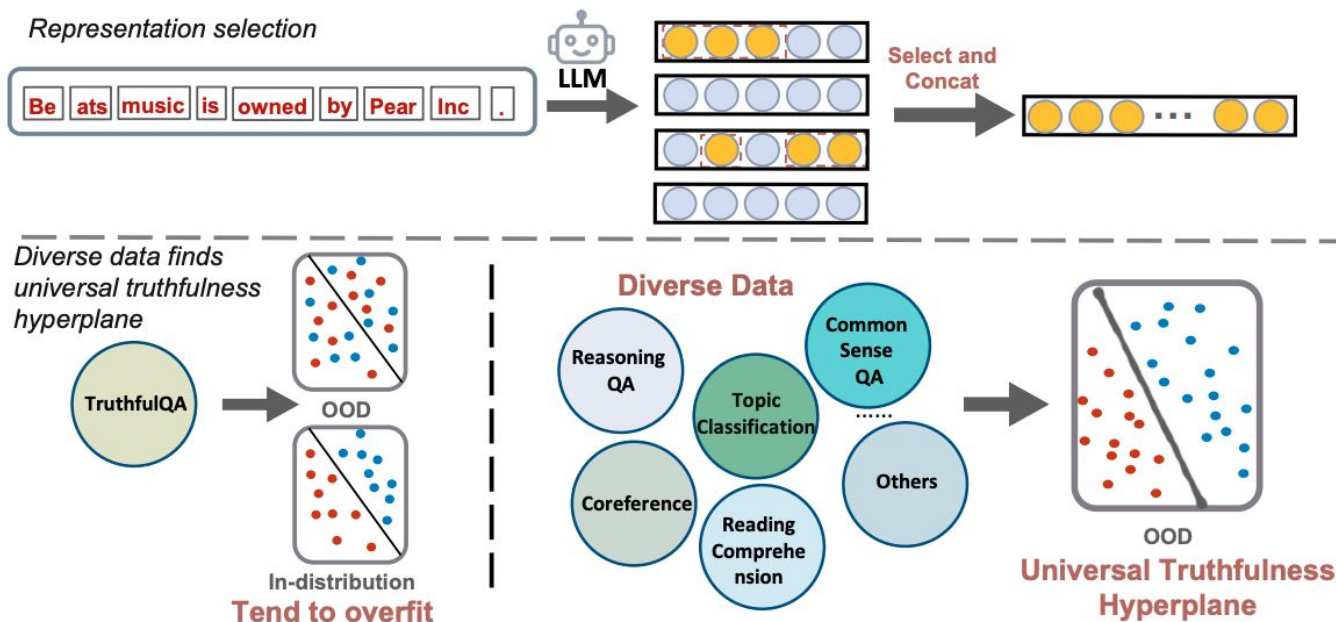


Figure 1: **Top:** we extract representations from the last token of the input sequence, then specific locations of the hidden states inside the LLM are selected and concatenated as input to train the probe. **Bottom:** Previous works mainly train the linear probe on one dataset which tends to overfit spurious features. Our work utilizes diverse datasets to examine whether a universal truthfulness hyperplane exists that can generalize to out-of-domain data.

Knowledge Boundary Identification – Summary

Uncertainty Estimation

- ❑ Uncertainty Decomposition
- ❑ Token Probability-based
- ❑ Semantic-based
- ❑ Conformal Prediction

Confidence Calibration

- ❑ Prompt-based Calibration
- ❑ Fine-tuning for Calibration

Internal State Probing

Identification approaches should be designed for different knowledge boundaries, suiting different mitigation approaches.